

Minimax Bounds Algorithms for MAB

Kiarash Gheisari, Arya Mokhtari, Mohammad Moradi

February 22, 2026

Multi-Armed Bandit Problem

- A decision maker repeatedly chooses among K actions (arms)
- Each arm gives a reward
- Rewards are unknown
- Goal: maximize cumulative reward over time

Key challenge:

Exploration vs Exploitation

- Explore: try different arms to learn their rewards
- Exploit: play the arm that looks best so far

Why "Multi-Armed Bandit" ?

Why "Multi-Armed Bandit"?



Each arm = one slot machine.

- ▶ Each machine has unknown payout
- ▶ You want to earn as much money as possible
- ▶ You must learn while playing

- 1 -

Each arm = one slot machine.

- Each machine has unknown payout
- You want to earn as much money as possible
- You must learn while playing

Adversarial Multi-Armed Bandit

Rewards are not generated from fixed distributions.

- Environment can choose rewards arbitrarily
- Rewards may depend on past actions
- No statistical assumptions

Interpretation:

- Environment may be hostile or strategic
- Past rewards may not predict future rewards

Goal:

Perform well against the worst possible sequence

Stochastic Multi-Armed Bandit

Each arm has a fixed reward distribution.

- Arm i has distribution ν_i
- Rewards are i.i.d. over time
- Mean reward of arm i is μ_i

Interpretation:

- The world is stable
- Past observations help predict future rewards

Goal:

Identify and play the best arm (highest mean)

Stochastic vs Adversarial Bandits

Stochastic	Adversarial
Fixed distributions	Arbitrary rewards
i.i.d. rewards	Possibly changing rewards
Learn the best arm	Compete with best fixed arm
Statistical inference	Worst-case guarantees

Adversarial model is strictly more general.

Bandit Interaction Protocol

For $t = 1, \dots, n$:

- 1 Learner chooses arm $I_t \in \{1, \dots, K\}$
- 2 Environment generates reward vector

$$g_t = (g_{1,t}, \dots, g_{K,t})$$

- 3 Learner observes only the reward of chosen arm:

$$g_{I_t,t}$$

This is called **partial feedback**.

A policy determines how to choose arms.

- Uses past actions and observed rewards
- Outputs probability distribution over arms

$$p_t = (p_{1,t}, \dots, p_{K,t})$$

Arm I_t is sampled from p_t .

Objective

Total reward of the learner:

$$\sum_{t=1}^n g_{I_t, t}$$

But how do we measure performance?

We compare with the best fixed arm in hindsight.

Regret compares learner to the best arm:

$$R_n = \max_{i=1,\dots,K} \mathbb{E} \left[\sum_{t=1}^n g_{i,t} - \sum_{t=1}^n g_{I_t,t} \right]$$

Interpretation:

- First term: reward of best fixed arm
- Second term: reward of learner

Goal:

R_n should grow as slowly as possible

Why Regret?

We cannot expect to always play the best arm.
But we want:

$$\frac{R_n}{n} \rightarrow 0$$

Meaning:

- Average performance approaches optimal
- Learner eventually behaves optimally

Why EXP3?

- No stochastic assumptions
- Rewards may be adversarial
- Deterministic algorithms fail

Goal:

$$R_n = \max_i \mathbb{E} \left[\sum_{t=1}^n g_{i,t} - g_{I_t,t} \right]$$

EXP3: Step 1 — Initialization

Initialize weights:

$$w_{i,1} = 1 \quad \forall i$$

At time t define probabilities:

$$p_{i,t} = \frac{w_{i,t}}{\sum_{j=1}^K w_{j,t}}$$

Sample arm $I_t \sim p_t$.

EXP3: Step 2 — Reward Estimation

Observe reward $g_{I_t,t} \in [0, 1]$.

Define unbiased estimator:

$$\hat{g}_{i,t} = \begin{cases} \frac{g_{i,t}}{p_{i,t}} & i = I_t \\ 0 & \text{otherwise} \end{cases}$$

- Corrects for partial feedback
- High variance

EXP3: Step 3 — Update

Update weights multiplicatively:

$$w_{i,t+1} = w_{i,t} \exp(\eta \hat{g}_{i,t})$$

- Arms with high estimated gain get higher weight
- Randomization ensures robustness

EXP3: Regret Bound

With $\eta \asymp \sqrt{\frac{\log K}{nK}}$:

$$R_n \leq C\sqrt{nK \log K}$$

- Optimal up to $\sqrt{\log K}$
- Gap motivates INF

Minimax Regret: What Is Possible?

Question:

What is the smallest regret any algorithm can guarantee?

We study the **minimax regret**:

$$\inf_{\pi} \sup_{\mu} R_n(\pi, \mu)$$

Main result (stochastic case):

$$\inf_{\pi} \sup_{\mu} R_n(\pi, \mu) \geq c\sqrt{nK}$$

- Holds for *any* policy
- Shows a fundamental difficulty of learning

Lower Bound: High-Level Strategy

Key idea: **confuse the learner**.

- Construct two bandit instances that are:
 - statistically hard to distinguish
 - have different optimal arms
- Any algorithm must fail on at least one

Lower Bound: High-Level Strategy

Key idea: **confuse the learner**.

- Construct two bandit instances that are:
 - statistically hard to distinguish
 - have different optimal arms
- Any algorithm must fail on at least one

This forces:

large regret in at least one instance

Lower Bound Construction

Two stochastic bandits:

Instance ν :

$$\mu_1 = \Delta, \quad \mu_2 = \cdots = \mu_K = 0$$

Instance ν' :

$$\mu_i = 2\Delta, \quad \mu_1 = \Delta, \quad \text{others } 0$$

- Only *one arm differs*
- Rewards are Gaussian

Why This Works

- Learner must explore all arms
- Some arm is played at most $\frac{n}{K}$ times
- That arm may secretly be optimal

Why This Works

- Learner must explore all arms
- Some arm is played at most $\frac{n}{K}$ times
- That arm may secretly be optimal

Tradeoff:

- Explore $\Rightarrow$ regret in ν
- Exploit $\Rightarrow$ regret in ν'

No escape.

Role of KL Divergence

- Learner sees only partial feedback
- Observations under ν and ν' are close

Key inequality:

$$P_\nu(A) + P_{\nu'}(A^c) \geq \frac{1}{2} e^{-D(P_\nu \| P_{\nu'})}$$

KL divergence controls distinguishability.

Lower Bound Takeaway

Optimizing the gap Δ yields:

$$R_n \geq c\sqrt{nK}$$

- No algorithm can beat this rate
- This is the benchmark for all upper bounds

INF uses a function $\psi : \mathbb{R} \rightarrow \mathbb{R}_+$ such that:

- ψ is increasing and C^1
- ψ'/ψ is nondecreasing
- $\lim_{u \rightarrow -\infty} \psi(u) < 1/K$
- $\lim_{u \rightarrow +\infty} \psi(u) \geq 1$

Normalization (Lemma 3)

For any vector $x \in \mathbb{R}^K$, there exists a unique $C(x)$ such that:

$$\sum_{i=1}^K \psi(x_i - C(x)) = 1$$

Moreover:

$$\max_i x_i - C(x) \leq \max_i x_i - \psi^{-1}(1/K)$$

Proof uses the Mean Value Theorem.

INF Algorithm

Maintain estimated cumulative gains $\hat{G}_{i,t}$.

Define probabilities:

$$p_{i,t} = \psi(\hat{G}_{i,t} - C_t) \quad \text{with} \quad \sum_i p_{i,t} = 1$$

Sample $I_t \sim p_t$.

INF: Gain Estimation

Define:

$$\hat{g}_{i,t} = \begin{cases} \frac{g_{i,t}}{p_{i,t}} & i = I_t \\ 0 & \text{otherwise} \end{cases}$$

Update:

$$\hat{G}_{i,t+1} = \hat{G}_{i,t} + \hat{g}_{i,t}$$

Theorem 4: INF Optimality

Theorem 4 For any real $q > 1$, the Implicitly Normalized Forecaster with

$$\psi(x) = \frac{1}{K} \left(\frac{9\sqrt{qnK}}{-x} \right)^q + \frac{q^{q/(2q-2)}}{\sqrt{nK}}$$

satisfies

$$\sup R_n \leq \frac{37}{1 - 1/q} \sqrt{qnK},$$

where the supremum is taken over all adversarial environments with rewards in $[0, 1]$.

- Achieves optimal $\sqrt{nK}$ rate
- Removes $\sqrt{\log K}$ gap from EXP3
- No explicit exploration

Stochastic Setting

- K arms with unknown means μ_i
- Rewards are i.i.d. and bounded in $[0, 1]$
- Gaps: $\Delta_i = \mu^* - \mu_i$

Regret decomposition:

$$R_n = \sum_{i: \Delta_i > 0} \Delta_i \mathbb{E}[T_i(n)]$$

UCB1 Algorithm

For each arm i at time t , define

$$\text{UCB}_i(t) = \hat{\mu}_i(t) + \sqrt{\frac{2 \log t}{T_i(t)}}$$

At each round:

- Play arm with largest UCB
- Optimism enforces exploration

Theorem (UCB1)

For all $n \geq 1$,

$$R_n \leq 10 \left(\sum_{i: \Delta_i > 0} \frac{1}{\Delta_i} \right) \log n$$

- Instance-dependent logarithmic regret
- Matches lower bounds up to constants

UCB1: Key Proof Ideas

- Concentration: empirical mean stays near μ_i
- A suboptimal arm i is played only when:
 - its estimate is overly optimistic, or
 - the optimal arm is overly pessimistic
- Both events are exponentially unlikely

This yields:

$$\mathbb{E}[T_i(n)] = O\left(\frac{\log n}{\Delta_i^2}\right)$$

Worst-Case Regret of UCB1

Uniform worst case over all instances:

$$R_n = \Theta(\sqrt{nK \log n})$$

- Extra $\sqrt{\log n}$ factor
- Does not match minimax $\sqrt{nK}$

Alternative Regret Bound Formulation

$$R_n \leq \max_{t_i \geq 0, \sum_i t_i = n} \sum_{i: \Delta_i > 0} \min \left(\frac{10}{\Delta_i} \log n, t_i \Delta_i \right)$$

- Shows the tradeoff in allocating pulls
- Each suboptimal arm contributes at most $\frac{10}{\Delta_i} \log n$ regret
- But total pulls are limited to n

The Stochastic Case: Motivation

In the deterministic case $g_{i,t} = 1$ if $i = 1$, 0 otherwise:

- INF policies and EXP3 have expected regret $\geq \sqrt{nK}$
- But we want smaller regret in stochastic settings
- Ideally: $\sqrt{nK}$ worst-case, but much better when gaps are large

Let $\Delta_i = \max_j \mu_j - \mu_i$ measure suboptimality.

MOSS Algorithm Definition

MOSS (Minimax Optimal Strategy in the Stochastic case):

For arm i , after s draws:

$$B_{i,s} = \hat{X}_{i,s} + \sqrt{\frac{\max(\log(\frac{n}{Ks}), 0)}{s}}$$

with $B_{i,0} = +\infty$.

At time t , draw arm maximizing $B_{i,T_i(t-1)}$.

- $\hat{X}_{i,s}$: empirical mean after s draws
- $T_i(t)$: number of draws of arm i up to time t

MOSS: Key Intuition

- Similar to UCB1 but with $\log(n/(Ks))$ instead of $\log t$
- Bonus decreases as arm gets more pulls
- Arms drawn more than n/K times use just empirical mean
- Automatically balances exploration and exploitation

Theorem 5: MOSS Minimax Bound

Theorem 5 MOSS satisfies:

$$\sup R_n \leq 49\sqrt{nK}$$

where the supremum is taken over all K -tuple of probability distributions on $[0, 1]$.

- Achieves optimal $\sqrt{nK}$ rate
- Matches lower bound up to constant factor
- Constant 49 can be improved to 5.7 with careful tuning

MOSS Proof: Decoupling Trick

Key idea: separate regret into manageable terms.

Define thresholds $z_k = \mu_1 - \frac{\Delta_k}{2}$ (midpoint between best and arm k).

Define $Z = \min_{1 \leq s \leq n} B_{1,s}$ (lowest index of best arm).

Define $\tau_k = \min\{t : B_{k,t} < z_k\}$ (when arm k 's index falls below threshold).

MOSS Proof: Crucial Inclusion

Key observation:

$$\{Z \geq z_k\} \subseteq \{T_k(n) \leq \tau_k\}$$

- If best arm's index stays above z_k
- Then once arm k 's index drops below z_k , it can't be played again
- So k cannot be pulled more than τ_k times

This decouples the arms!

MOSS Proof: Regret Decomposition

After algebra, we get:

$$R_n \leq 2n\Delta_{k_0} + \sum_{k > k_0} \Delta_k \mathbb{E}[\tau_k] + n \sum_{k > k_0} \mathbb{P}(Z < z_k)(\Delta_k - \Delta_{k-1})$$

Three terms:

- ① $2n\Delta_{k_0}$: crude bound for "small-gap" arms
- ② $\sum \Delta_k \mathbb{E}\tau_k$: time to rule out each suboptimal arm
- ③ $n \sum \mathbb{P}(Z < z_k)(\Delta_k - \Delta_{k-1})$: cost when best arm dips

MOSS Proof: Choosing the Split

Pick threshold:

$$\delta_0 = e^{1/16} \sqrt{\frac{K}{n}}$$

Choose k_0 so that:

$$\Delta_{k_0} \leq \delta_0 < \Delta_{k_0+1}$$

- Arms $1, \dots, k_0$: tiny gaps (hard)
- Arms $k_0 + 1, \dots, K$: gaps at least δ_0 (easier)

MOSS Proof: Bounding $\mathbb{E}\tau_k$

$$\mathbb{E}[\tau_k] = \sum_{\ell \geq 0} \mathbb{P}(\tau_k > \ell) = \sum_{\ell \geq 0} \mathbb{P}(\forall t \leq \ell, B_{k,t} \geq z_k)$$

For $\ell \geq \ell_0 \asymp \frac{\log(n/(K\Delta_k^2))}{\Delta_k^2}$:

$$\mathbb{P}(B_{k,\ell} \geq z_k) \leq \mathbb{P}(\hat{\mu}_{k,\ell} - \mu_k \geq c\Delta_k) \leq e^{-2\ell c^2 \Delta_k^2}$$

Summing yields:

$$\Delta_k \mathbb{E}\tau_k \lesssim 1 + \frac{\log(n/(K\Delta_k^2))}{\Delta_k}$$

Using $\Delta_k \geq \delta_0$ and maximization:

$$\sum_{k > k_0} \Delta_k \mathbb{E}\tau_k \leq K + 28.3\sqrt{nK}$$

MOSS Proof: Bounding $\mathbb{P}(Z < z_k)$

$Z < z_k = \mu_1 - \Delta_k/2$ means: at some sample size s , empirical mean of arm 1 dipped unusually low.

Translate to partial sums:

$$\sum_{t=1}^s (\mu_1 - X_t) \text{ exceeding a boundary}$$

Control using:

- Maximal inequality (Hoeffding for partial sums)
- Peeling argument (split s into geometric ranges)

MOSS Proof: Peeling Computation

After computations:

$$n \sum_{k > k_0} \mathbb{P}(Z < z_k)(\Delta_k - \Delta_{k-1}) \leq n(\delta_0 - \Delta_{k_0}) + 17.8\sqrt{nK}$$

- "Best arm dips" events do happen
- But total contribution is still $O(\sqrt{nK})$

MOSS Proof: Final Combination

Plugging bounds into decomposition:

$$R_n \leq 48.3\sqrt{nK} + K$$

Two regimes:

- If K large compared to n : trivially $R_n \leq n \leq 49\sqrt{nK}$
- Otherwise $K \leq \frac{1}{49}\sqrt{nK}$, so $48.3\sqrt{nK} + K \leq 49\sqrt{nK}$

Thus:

$$R_n \leq 49\sqrt{nK}$$

Comparison of Algorithms

Algorithm	Regret Bound
UCB1	$\min \left(\sqrt{nK \log n}, \sum_{i: \Delta_i > 0} \frac{\log n}{\Delta_i} \right)$
MOSS	$\min \left(\sqrt{nK}, \sum_{i: \Delta_i > 0} \frac{K \log(2 + n\Delta_i^2/K)}{\Delta_i} \right)$
EXP3	$\sqrt{nK \log K}$
INF	$\sqrt{nK}$

- MOSS achieves best of both worlds
- INF optimal in adversarial setting

Summary: Key Takeaways

- Multi-armed bandits model exploration-exploitation tradeoff
- Two main settings: stochastic and adversarial
- Lower bound: $\Omega(\sqrt{nK})$ regret is unavoidable
- UCB1: $\tilde{O}(\log n)$ instance-dependent, $O(\sqrt{nK \log n})$ worst-case
- EXP3: $O(\sqrt{nK \log K})$ in adversarial setting
- INF: Optimal $O(\sqrt{nK})$ in adversarial setting
- MOSS: Optimal $O(\sqrt{nK})$ in stochastic setting with better instance-dependent bounds

Potential Improvements to EXP3

Issue: High Variance Estimator

EXP3 uses unbiased importance sampling estimator:

$$\hat{g}_{i,t} = \frac{g_{i,t}}{p_{i,t}} \mathbb{1}_{\{I_t=i\}}$$

- Variance can be huge when $p_{i,t}$ is small
- Can destabilize learning and slow convergence

Two Possible Improvements:

Improvement 1: Biased Low-Variance Estimators

Replace unbiased estimator with biased but lower variance alternative:

Weighted Importance Sampling:

$$\hat{g}_{i,t} = \frac{w_{i,t} g_{i,t}}{\sum_j w_{j,t} p_{j,t}} \mathbb{1}_{\{I_t=i\}}$$

- Trade off bias for reduced variance
- Can lead to faster empirical convergence
- Similar to difference between ordinary and weighted importance sampling in RL

Open question: Can we bound the regret with biased estimators?

Improvement 2: Baseline Subtraction

Add a baseline b_t to reduce variance without introducing bias:

EXP3 with Baseline:

$$\hat{g}_{i,t} = \frac{g_{i,t} - b_t}{p_{i,t}} \mathbb{1}_{\{I_t=i\}} + b_t$$

- Baseline b_t can be any value known before choosing action
- Common choice: $b_t = \sum_i p_{i,t} \hat{g}_{i,t}$ (weighted average)
- Preserves unbiasedness: $\mathbb{E}[\hat{g}_{i,t}] = g_{i,t}$
- Reduces variance when baseline correlates with rewards

Effect: Smoother updates, better empirical performance

Open Questions

- Can we achieve both optimal stochastic and adversarial guarantees simultaneously?
- What about contextual bandits?
- How to handle very large action spaces?
- Best constants in minimax bounds?

References

-  Auer, P., Cesa-Bianchi, N., & Fischer, P. (2002). Finite-time analysis of the multiarmed bandit problem.
-  Audibert, J.-Y., & Bubeck, S. (2009). Minimax policies for adversarial and stochastic bandits.
-  Bubeck, S., & Cesa-Bianchi, N. (2012). Regret analysis of stochastic and nonstochastic multi-armed bandit problems.